

Multicoretests – Parallel Testing Libraries for OCaml 5.0

Jan Midtgaard

Olivier Nicole

Nicolas Osborne

Tarides

1 Introduction

Parallel and concurrent code is notoriously hard to test because of the involved non-determinism, yet it is facing OCaml programmers with the coming OCaml 5.0 multicore release. We present two related testing libraries to improve upon the situation:

- `Lin` – a library to test for linearizability
- `STM` – a state-machine testing library

Both libraries build on `QCheck` [18], a black-box, property-based testing library in the style of `QuickCheck` [5]. The two libraries represent different trade-offs between required user effort and provided guarantees and thereby supplement each other.

In this document we will use OCaml’s `Hashtbl` module as a running example.

2 The `Lin` library

The `Lin` library performs a sequence of random operations in parallel, records the results, and checks whether the observed results are linearizable by reconciling them with a sequential execution. The library offers an embedded, combinator DSL to describe signatures succinctly. As an example, the required specification to test (parts of) the `Hashtbl` module is given in fig. 1.

The first line indicates the type of the system under test (SUT). In the above case we intend to test `Hashtbl`s with char keys and int values. The bindings `init` and `cleanup` allow for setting up and tearing down the SUT. The `api` then contains a list of type signature descriptions using combinators in the style of `Ctypes` [23]. Different combinators `unit`, `bool`, `int`, `list`, `option`, `returning`, `returning_or_exc`, ... allow for a concise type signature description.

From the signature description the `Lin` library will iterate a number of test instances. Each test instance consists of a “sequential prefix” of calls to the specified operations, followed by a spawn of two parallel `Domains` that each call a sequence of operations.

For each test instance `Lin` chooses the individual operations arbitrarily and records the result received from each operation. The framework will then perform a search for a sequential interleaving of the same calls, and succeed if it finds one. Since `Hashtbl`s are not safe for parallelism, the output produces the following:

Results incompatible with sequential execution

```

      |
      | Hashtbl.add t '@' 4 : ()
      |
      |-----|
      |         |
Hashtbl.add t '.' 3 : ()  Hashtbl.clear t : ()
                          |
                          | Hashtbl.length t : 2
```

This describes that in one parallel execution, `Lin` received the response 2 from `Hashtbl.length`, despite having just executed `Hashtbl.clear`. In this case, it is not possible to interleave `Hashtbl.add t '.' 3` with these two calls to explain this observed behaviour.

Underneath the hood, `Lin` does its best to schedule the two parallel `Domains` on top of each other. It also repeats each test instance, to increase the chance of triggering an error, and it fails if just one of the repetitions fail to find a sequential interleaving. Finally, upon finding an error it reduces the involved operation sequences to a local minimum, which is what is printed above.

`Lin` is phrased as an OCaml functor, `Lin_api.Make`. The module resulting from `Lin_api.Make(HashtblSig)` contains a binding `lin_test` that can perform the above linearizability test over `Domains`, the basic unit of parallelism coming in OCaml 5.0. An alternative `Lin` mode works over `Thread` for testing concurrent but non-overlapping executions. This mode thus mimicks the above functionality by replacing `Domain.spawn` and `Domain.join` with `Thread.create` and `Thread.join`, respectively.

3 The `STM` library

Like `Lin` the `STM` library also performs a sequence of random operations in parallel and records the results. In contrast to `Lin`, `STM` then checks whether the observed results are linearizable by reconciling them with a sequential execution of a model description.

The model expresses the intended meaning of each tested operation. As such, the required `STM` user input is longer compared to that of `Lin`. The corresponding code to describe a `Hashtbl` test using `STM` is given in fig. 2.

Again this requires a description of the system under test, `sut`. In addition `STM` requires a type `cmd` for describing the tested operations. The hooks `init_sut` and `cleanup` match `init` and `cleanup` from `Lin`, respectively.

```

module HashtblSig =
struct
  type t = (char, int) Hashtbl.t
  let init () = Hashtbl.create ~random:false 42
  let cleanup _ = ()

  open Lin_api
  let a,b = char_printable,nat_small
  let api =
    [ val_ "Hashtbl.clear"      Hashtbl.clear      (t @-> returning unit);
      val_ "Hashtbl.add"        Hashtbl.add        (t @-> a @-> b @-> returning unit);
      val_ "Hashtbl.remove"     Hashtbl.remove     (t @-> a @-> returning unit);
      val_ "Hashtbl.find"       Hashtbl.find       (t @-> a @-> returning_or_exc b);
      val_ "Hashtbl.replace"    Hashtbl.replace    (t @-> a @-> b @-> returning unit);
      val_ "Hashtbl.mem"        Hashtbl.mem        (t @-> a @-> returning bool);
      val_ "Hashtbl.length"     Hashtbl.length     (t @-> returning int); ]
end

```

Figure 1: Specification of selected Hashtbl functions for testing using Lin.

A distinguishing feature is type `state = (char * int) list` describing with a pure association list the internal state of a hashtable. `next_state` is a simple state transition function describing how the state changes across each `cmd`. For example, `Add (k,v)` appends the key-value pair onto the association list.

`arb_cmd` is a generator of `cmds`, taking `state` as a parameter. This allows for state-dependent `cmd` generation, which we use to increase the chance of producing a `Remove 'c'`, `Find 'c'`, ... following an `Add 'c'`. Internally `arb_cmd` uses combinators `Gen.return`, `Gen.map`, and `Gen.map2` from `QCheck` to generate one of 7 different operations. For example, `Gen.map (fun k -> Mem k) char` creates a `Mem` command with the result obtained from the `char` generator. `arb_cmd` further uses a derived printer `show_cmd` to be able to print counterexamples.

`run` executes the tested `cmd` over the SUT and wraps the result up in a result type `res` offered by `STM`. Combinators `unit`, `bool`, `int`, ... allow to annotate the result with the expected type. `postcond` then expresses a post-condition by matching the received `res`, for a given `cmd` with the corresponding answer from the model description. For example, this compares the Boolean result `r` from `Hashtbl.mem` with the result from `List.mem_assoc`. Similarly `precond` expresses a `cmd` pre-condition.

`STM` is also phrased as an OCaml functor. The module resulting from `STM.Make(HashtblModel)` thus includes a binding `agree_test` for running sequential tests comparing the SUT behaviour to the given model. Another binding `agree_test_par` instead runs parallel tests that make a similar comparison over a sequential prefix and two parallel Domains, this time also searching for a sequential interleaving of `cmds`. For example, one execution of `agree_test_par` produced the following output. Note how no interleaving of `Remove` from the first parallel `cmd` sequence can make the association list model return `-1` from `Length`:

Results incompatible with linearized model

```

              |
          (Add ('1', 5)) : ()
              |
-----
|               |
(Remove '1') : ()      Clear : ()
                        Length : -1

```

4 Status

Both libraries are open source and available for download on GitHub from <https://github.com/jmid/multicoretests>. As the APIs are still unstable and under development, we have not made a public release yet. Interested users can nevertheless easily install the libraries with `opam`.

During development we have used examples such as `Hashtbl` to confirm that the approach indeed works as intended. The behaviour is continuously confirmed by running GitHub Actions of the latest trunk compiler. As further testament to the usability of the approach, we have used the libraries to test parts of OCaml's `Stdlib`, as well as the `Domainlib` and `lockfree` libraries. In doing so, we have been able to find and report a number of issues which have either already been fixed or have fixes underway:

- `In_channel` and `Out_channel` unsafety [1, 3]
- MacOSX crash [21]
- Buffer unsafety [22, 15]

5 Related Work

`QuickCheck` [5] originally introduced property-based testing within functional programming with combinator-based generators, properties, and test-case reduction. It has since been ported to over 30 other programming languages, including Quviq `QuickCheck` [19]—a commercial port to Erlang.

```

module HashtblModel =
struct
  type sut = (char, int) Hashtbl.t
  type state = (char * int) list
  type cmd =
    | Clear
    | Add of char * int
    | Remove of char
    | Find of char
    | Replace of char * int
    | Mem of char
    | Length [@@deriving show { with_path = false }]

  let init_sut () = Hashtbl.create ~random:false 42
  let cleanup (_:sut) = ()

  let arb_cmd (s:state) =
    let char =
      if s = []
      then Gen.printable
      else Gen.(oneof [oneofl (List.map fst s);
                        printable]) in
    let int = Gen.nat in
    QCheck.make ~print:show_cmd
      (Gen.oneof
        [Gen.return Clear;
         Gen.map2 (fun k v -> Add (k,v)) char int;
         Gen.map (fun k -> Remove k) char;
         Gen.map (fun k -> Find k) char;
         Gen.map2 (fun k v -> Replace (k,v)) char int;
         Gen.map (fun k -> Mem k) char;
         Gen.return Length;
        ])

  let next_state (c:cmd) (s:state) = match c with
    | Clear      -> []
    | Add (k,v)  -> (k,v)::s
    | Remove k   -> List.remove_assoc k s
    | Find _     -> s
    | Replace (k,v) -> (k,v)::(List.remove_assoc k s)
    | Mem _      -> s
    | Length     -> s

  let run (c:cmd) (h:sut) = match c with
    | Clear      -> Res (unit, Hashtbl.clear h)
    | Add (k,v)  -> Res (unit, Hashtbl.add h k v)
    | Remove k   -> Res (unit, Hashtbl.remove h k)
    | Find k     -> Res (result int exn,
                        protect (Hashtbl.find h) k)
    | Replace (k,v) -> Res (unit, Hashtbl.replace h k v)
    | Mem k      -> Res (bool, Hashtbl.mem h k)
    | Length     -> Res (int, Hashtbl.length h)

  let init_state = []

  let precondition (_:cmd) (_:state) = true
  let postcondition (c:cmd) (s:state) (res:res) =
    match c,res with
    | Clear,      Res ((Unit,_),_)
    | Add (_,_),  Res ((Unit,_),_)
    | Remove _,   Res ((Unit,_),_) -> true
    | Find k,     Res ((Result (Int,Exn),_),r) ->
      r = (try Ok (List.assoc k s)
            with Not_found -> Error Not_found)
    | Replace (_,_), Res ((Unit,_),_) -> true
    | Mem k,       Res ((Bool,_),r) -> r = List.mem_assoc k s
    | Length,      Res ((Int,_),r) -> r = List.length s
    | _ -> false
end

```

Figure 2: Description of a Hashtbl test using STM.

Model-based testing was initially suggested as a method for testing monadic code with Haskell’s QuickCheck [6]. An explicit framework was later proposed in the GAST property-based testing library for Clean [10]. The commercial Quviq QuickCheck [19] was later extended with a state-machine model framework for testing stateful systems [2]. This approach was extended further to test parallel code for data races [7]. This general approach for parallel testing has since been adopted in other ports, such as Erlang’s open source Proper [14], Haskell Hedgehog [9], ScalaCheck [20], and Kotlin’s propCheck [17]. STM continues this adoption tradition. qcstm [11] is a previous OCaml adoption, also building on QCheck. It was missing the ability to perform parallel testing though. STM seeks to remedy this limitation.

Crowbar [8] is another QuickCheck-style testing framework with combinator-based generators. In contrast to QuickCheck, it utilizes AFL-based coverage guidance to effectively guide the generated input towards unvisited parts of the SUT. Crowbar does not come with a state-machine framework. Monolith [16] is a model-based testing framework also building on AFL-based coverage guidance. In contrast to STM, Monolith’s models are oracle implementations

with operations matching the type signatures of the tested operations. Neither Crowbar nor Monolith come with skeletons to perform parallel or concurrent testing. Furthermore the AFL-based coverage-guidance underlying both Crowbar and Monolith works best for deterministic, sequential code.

ParaFuzz [13] is another approach to fuzz test multicore OCaml programs. It simulates parallelism in OCaml through concurrency, enabling scheduling order to be controlled by AFL, which helps to trigger and find scheduling-dependent bugs. A caveat is that ParaFuzz assumes data race freedom.

Ortac can extract Monolith-based tests from a formal specification written in Gospel, a specification language for OCaml [12]. Gospel specifications include models, preconditions, and post-conditions close to those of STM. The extracted tests however inherit Monolith’s and AFL’s focus on sequential code.

ArtiCheck [4] tests random combinations of OCaml calls from type signature descriptions, similarly to Lin. Whereas Lin and STM target impure interfaces, ArtiCheck targets persistent (pure) interfaces. ArtiCheck furthermore targets sequential rather than parallel or concurrent tests.

6 Conclusion

We have presented two libraries, `Lin` and `STM` for testing parallel and concurrent code for OCaml 5.0. Despite still being under development, we believe both libraries could be helpful to developers of OCaml 5.0 programs.

References

- [1] *Add (Failing) {In,Out_channel Linearization Tests*. URL: <https://github.com/jmid/multicoretests/pull/13>.
- [2] Thomas Arts et al. “Testing Telecoms Software with Quviq QuickCheck”. In: *Proceedings of the 2006 ACM SIGPLAN Workshop on Erlang (Erlang 2006)*. 2006, pp. 2–10.
- [3] *Audit Stdlib for Mutable State (Comment)*. URL: <https://github.com/ocaml/ocaml/issues/10960#issuecomment-1087660763>.
- [4] Thomas Braibant, Jonathan Protzenko, and Gabriel Scherer. “Well-Typed Generic Smart Fuzzing for APIs”. In: *ML Family Workshop (ML 2014)*. 2014. URL: <https://hal.inria.fr/hal-01094006>.
- [5] Koen Claessen and John Hughes. “QuickCheck: A Lightweight Tool for Random Testing of Haskell Programs”. In: *Proceedings of the Fifth ACM SIGPLAN International Conference on Functional Programming (ICFP 2000)*. 2000, pp. 268–279.
- [6] Koen Claessen and John Hughes. “Testing Monadic Code with QuickCheck”. In: *Proceedings of the 2002 ACM SIGPLAN Workshop on Haskell (Haskell 2002)*. 2002, pp. 65–77.
- [7] Koen Claessen et al. “Finding Race Conditions in Erlang with QuickCheck and PULSE”. In: *Proceeding of the 14th ACM SIGPLAN International Conference on Functional Programming (ICFP 2009)*. 2009, p. 12.
- [8] Stephen Dolan and Mindy Preston. “Testing with Crowbar”. In: *OCaml Users and Developers Workshop*. 2017.
- [9] *Hedgehog*. URL: <https://github.com/hedgehogqa/haskell-hedgehog>.
- [10] Pieter W. M. Koopman and Rinus Plasmeijer. “Testing Reactive Systems with GAST”. In: *Revised Selected Papers from the Fourth Symposium on Trends in Functional Programming (TFP 2003)*. Vol. 4. Trends in Functional Programming. 2003, pp. 111–129.
- [11] Jan Midtgaard. “A Simple State-Machine Framework for Property-Based Testing in OCaml”. In: *OCaml Users and Developers Workshop*. 2020.
- [12] Nicolas Osborne and Clément Pascutto. “Leveraging Formal Specifications to Generate Fuzzing Suites”. In: *OCaml Users and Developers Workshop*. 2021. URL: <https://hal.inria.fr/hal-03328646>.
- [13] Sumit Padhiyar, Adharsh Kamath, and KC Sivaramakrishnan. “Parafuzz: Coverage-guided Property Fuzzing for Multicore OCaml Programs”. In: *OCaml Users and Developers Workshop*. 2021.
- [14] Manolis Papadakis and Konstantinos Sagonas. “A PropEr Integration of Types and Function Specifications with Property-Based Testing”. In: *Proceedings of the 2011 ACM SIGPLAN Erlang Workshop*. 2011, pp. 39–50.
- [15] *Parallel Access to Buffer Can Trigger Segfaults*. URL: <https://github.com/ocaml/ocaml/issues/11279>.
- [16] François Pottier. “Strong Automated Testing of OCaml Libraries”. In: *Journées Francophones Des Langues Applicatifs (JFLA 2021)*. Feb. 2021.
- [17] *propCheck*. URL: <https://github.com/1Jajen1/propCheck>.
- [18] *QCheck*. URL: <https://github.com/c-cube/qcheck>.
- [19] *Quviq QuickCheck*. URL: <http://quviq.com/documentation/eqc/index.html>.
- [20] *ScalaCheck*. URL: <https://github.com/typelevel/scalacheck>.
- [21] *Segfault on MacOSX with Trunk*. URL: <https://github.com/ocaml/ocaml/issues/11226>.
- [22] *STM Clean-Up*. URL: <https://github.com/jmid/multicoretests/pull/63>.
- [23] Jeremy Yallop, David Sheets, and Anil Madhavapeddy. “A modular foreign function interface”. In: *Science of Computer Programming* 164 (2018), pp. 82–97. URL: <https://www.sciencedirect.com/science/article/pii/S0167642317300709>.